

Haripriya Varanasi | Data Analyst

Fairfax, VA | +1 (571) 278-3169 | haripriyavaranasi0106@gmail.com | LinkedIn | GitHub | Portfolio

SUMMARY

Data Analyst with **4+ years of progressive experience** turning operational, customer, and business data into reliable reporting and decision-ready insights. Hands-on with **SQL, Python, Power BI, Tableau, Excel, Databricks, and PySpark** across KPI reporting, customer analysis, data transformation, and quality validation. Experienced working with **AWS, ETL/ELT workflows, Snowflake, dbt, and statistical analysis**, with practical exposure to forecasting, classification, and segmentation using Scikit-learn. Combines strong analytical fundamentals with modern cloud data tools and **applied AI/ML capabilities** while keeping business analysis and measurable outcomes at the center of the work.

TECHNICAL SKILLS

Programming & Query: Python, SQL, R, PySpark

Business Intelligence & Visualization: Power BI, DAX, Power Query, Tableau, Excel, PivotTables, Plotly, Matplotlib, Streamlit

Data Analysis: Pandas, NumPy, Exploratory Data Analysis (EDA), Statistical Analysis, Trend Analysis, Variance Analysis, Cohort Analysis, Root-Cause Analysis, KPI Development

Data Engineering & Modeling: ETL/ELT, Data Pipelines, Data Transformation, Data Modeling, Data Warehousing, Data Quality, Data Validation, Apache Spark, dbt, DuckDB

Cloud & Data Platforms: Databricks, Snowflake, AWS (S3, EC2, Lambda, RDS, IAM)

Machine Learning: Scikit-learn, XGBoost, Random Forest, Logistic Regression, Regression, Classification, Clustering, K-Means, PCA, Forecasting, Model Evaluation

AI & LLM: RAG, FAISS, Sentence-Transformers, Ollama, Vector Search, AI-Assisted Analytics

Tools & Delivery: Git, GitHub, GitHub Actions, Jupyter Notebook, REST APIs, Agile/Scrum

WORK EXPERIENCE

Perficient | USA

Jan 2026 – Present

Data Analyst

- Automated monthly performance analysis with **Python, Pandas, and parameterized SQL**, cutting preparation time by **24%** and giving business users faster visibility into operational trends.
- Replaced manually consolidated KPI workbooks with interactive **Power BI** pages using DAX, Power Query, drill-through views, and exception indicators, reducing reconciliation effort by **18%**.
- Processed **1M+ records through Databricks and PySpark**, resolving schema differences and shaping raw feeds into analytical tables suitable for historical and cross-functional comparisons.
- Introduced **20+ automated controls** for nulls, duplicate keys, mapping failures, schema drift, and calculation exceptions, catching upstream issues before they affected published metrics.
- Used advanced **SQL—CTEs, window functions, CASE expressions, and subqueries**—to trace unusual metric movements to transaction-level activity and distinguish genuine changes from system anomalies.
- Translated stakeholder requests into **KPI definitions, calculation logic, filters, and visualization specifications**, contributing assigned enhancements while keeping measures consistent with agreed business rules.
- Diagnosed failed **AWS/Databricks pipeline** runs through record-count comparisons, transformation checks, and source-to-target mapping reviews, documenting root causes for coordination with engineering teams.
- Evaluated a targeted forecasting use case with **Scikit-learn**, comparing predictions against historical baselines and communicating model limitations alongside findings for business planning discussions.

Quantiphi | India

Apr 2022 – Jul 2024

Data Analyst

- Examined customer behavior and operational performance with **Python and SQL**, uncovering recurring patterns that shortened analytical turnaround by **21%** compared with spreadsheet-heavy investigation.
- Shifted recurring KPI consumption from static files to **Tableau and Power BI**, enabling self-service slicing by customer, period, and category while lowering report assembly effort **17%**.
- Segmented **500K+ customer records** by behavioral and operational characteristics, comparing cohorts and period-over-period movement to identify groups contributing disproportionately to performance changes.
- Created **15+ Python profiling checks** during ingestion to surface malformed fields, category shifts, reference mismatches, and completeness issues before analysts began downstream exploration.
- Reconciled source totals during an **ETL migration**, documenting transformation rules and metric definitions so legacy and replacement reports produced comparable business measures during transition.
- Turned ambiguous business questions into targeted analyses, combining **trend, variance, cohort, and root-cause analysis** to explain performance drivers through concise visual narratives.
- Explored classification and segmentation models in Scikit-learn for suitable customer use cases, comparing machine-learning outputs with descriptive baselines before recommending patterns for further evaluation.

innData Analytics Private Limited | India

Jun 2021 – Mar 2022

Associate Data Analyst

- Reduced recurring correction requests by **15%** through structured **Excel and SQL reconciliation**, correcting inconsistent values and maintaining cleaner inputs for scheduled business reports.
- Standardized frequently used spreadsheets with **PivotTables, lookup functions, formulas, and reusable templates**, cutting monthly preparation effort by approximately **12%** for senior analysts.
- Retrieved **100K+ records** from relational tables using filters, joins, CASE statements, GROUP BY, and aggregations to fulfill scheduled extracts and one-off information requests.

- Maintained accuracy across **10+ key reporting fields** by comparing identifiers, totals, dates, and categorical values against source records before completing assigned deliverables.
- Assisted with monthly **trend and variance reviews**, organizing supporting details and highlighting unusual movements for senior analysts to investigate before communicating conclusions to business users.
- Introduced small **Python/Pandas utilities** for repetitive file reshaping and cleanup, reducing dependence on manual spreadsheet steps while building practical automation experience.

EDUCATION

Master of Science in Data Analytics Engineering George Mason University Fairfax, VA	Aug 2024 – May 2026
Bachelor of Technology in Computer Science Engineering Andhra University Visakhapatnam, Andhra Pradesh, India	Aug 2020 – May 2024

PROJECTS

Consumer Complaint Decision-Support Platform Python, SQL, DuckDB, Scikit-learn, FAISS, Sentence-Transformers, Ollama, Streamlit • Engineered an analytical platform across 16M+ CFPB complaint records, restructuring the application and query layer to reduce page-load time from 25+ seconds to under 2 seconds. • Trained and evaluated Random Forest models for complaint outcomes and company benchmarking using chronological splits, ROC-AUC, Brier score, and calibration analysis rather than relying solely on accuracy. • Integrated citation-grounded RAG, natural-language-to-SQL, and assisted drafting with FAISS and local LLMs, adding retrieval thresholds, query safeguards, and output traceability checks. • Added a 36-test Pytest suite covering SQL-injection safeguards, model evaluation, analytical calculations, and AI-output validation before deploying a reduced public demonstration environment.	
Modern Analytics Engineering Warehouse Snowflake, dbt Core, SQL, GitHub Actions • Modeled 1M+ retail records in Snowflake and dbt, replacing 20+ ad-hoc analytical queries with six reusable marts structured for consistent downstream reporting. • Added dbt tests and GitHub Actions CI for transformation validation, reducing modeled processing time by 60–75% while improving repeatability across analytical workflows.	
US Air Quality & Community Health Risk Platform Python, AWS S3, REST APIs, GitHub Actions, Streamlit • Integrated EPA, CDC PLACES, and U.S. Census APIs into a standardized analytical dataset covering 3,000+ U.S. counties, combining environmental, health, and demographic indicators. • Automated ingestion and validation through Python, AWS S3, and GitHub Actions, then surfaced county-level trends and risk patterns through an interactive Streamlit application.	